

Parallel Programming in Rust



Lukáš Hozda

2025-04-10

Brains Systems s.r.o

Language Model & Fundamentals



- Multiple threads accessing shared data
- At least one thread is writing
- No synchronization between accesses
- Undefined behavior in unsafe code
- Impossible in safe Rust



- Race conditions: Logic-level timing issues
- Safe Rust allows race conditions
- Race conditions cannot violate memory safety alone
- Only unsafe code + race condition can break safety

Example: Safe Race Condition



```
use std::thread;
use std::sync::atomic::{AtomicUsize, Ordering};
use std::sync::Arc;

let data = vec![1, 2, 3, 4];
let idx = Arc::new(AtomicUsize::new(0));
let other_idx = idx.clone();

thread::spawn(move || {
    other_idx.fetch_add(10, Ordering::SeqCst);
});

// Race condition: May panic, but memory safe
println!("{}", data[idx.load(Ordering::SeqCst)]);
```

Example: Unsafe Race Condition



```
use std::thread;
use std::sync::atomic::{AtomicUsize, Ordering};
use std::sync::Arc;

let data = vec![1, 2, 3, 4];
let idx = Arc::new(AtomicUsize::new(0));
let other_idx = idx.clone();

thread::spawn(move || {
    other_idx.fetch_add(10, Ordering::SeqCst);
});

if idx.load(Ordering::SeqCst) < data.len() {
    unsafe { // UNSAFE: idx could change after bounds check
        println!("{}", data.get_unchecked(idx.load(Ordering::SeqCst)));
    }
}
```



- Single writer XOR multiple readers
- Enforced at compile time
- Extends to concurrent code
- Built around preventing data races



- Send: Safe to transfer between threads
- Sync: Safe to share reference between threads
- Automatically derived for most types
- Unsafe traits - incorrect impl causes UB
- T is Sync if and only if &T is Send



Not Send or Sync:

- Raw pointers
- Rc (refcount not threadsafe)
- Non-threadsafe cells (UnsafeCell is Send, but not Sync)

Send and Sync:

- Most standard collections
- Arc<T> where T: Send + Sync
- Atomic types



- OS-level threads
- Guaranteed preemptive multitasking
- Heavier than async tasks
- Can use all CPU cores

```
use std::thread;
```

```
let handle = thread::spawn(|| {  
    // New thread  
    println!("Hello from thread!");  
});
```

```
handle.join().unwrap(); // Wait for completion
```



- Closure must be 'static
- Use move keyword
- Values are transferred between threads

```
let v = vec![1, 2, 3];
```

```
let handle = thread::spawn(move || {  
    println!("Vector: {:?}", v);  
}); // v is moved into thread
```

```
// Cannot use v here anymore  
handle.join().unwrap();
```



- Allows borrowing from parent scope
- Guarantees threads finish before scope ends
- More flexible than regular threads

```
let mut v = vec![1, 2, 3];
thread::scope(|s| {
    s.spawn(|| {
        // Can borrow v here!
        println!("First: {:?}", &v);
    });
    s.spawn(|| {
        v.push(4); // But we cannot mutate here!
    });
});
v.push(5); // We can mutate here :)
```



- Arc = Atomic Reference Counting
- Thread-safe shared ownership
- Clone = increment reference count
- Drop = decrement reference count

```
use std::sync::Arc;
```

```
let data = Arc::new(vec![1, 2, 3]);
```

```
let data_clone = Arc::clone(&data);  
thread::spawn(move || {  
    println!("Thread sees: {:?}", *data_clone);  
});
```



- Rc = single-threaded reference counting
- Arc = atomic (thread-safe) reference counting
- Arc has small performance overhead
- Cannot convert between them

```
use std::rc::Rc;      // Wrong!  
use std::sync::Arc;    // Correct!
```

```
let bad = Rc::new(42);  
// thread::spawn(move || { // Compile error!  
//     println!("{}", bad);  
// });
```

```
let good = Arc::new(42);  
thread::spawn(move || { // Works!  
    println!("{}", good);  
});
```

Shared Mutable State



- Mutual exclusion
- Only one thread can access data at once
- Others must wait
- RAII-style locking

```
use std::sync::{Arc, Mutex};
```

```
let counter = Arc::new(Mutex::new(0));
```

```
let counter2 = Arc::clone(&counter);
```

```
thread::spawn(move || {  
    let mut num = counter2.lock().unwrap();  
    *num += 1;  
}); // lock automatically released here
```




- Smaller and faster than `std::sync::Mutex`
- No poisoning on panic
- More efficient spinning
- Platform-specific optimizations

```
use parking_lot::Mutex;
```

```
let mutex = Mutex::new(0);  
{  
    let mut guard = mutex.lock(); // No unwrap needed!  
    *guard += 1;  
} // lock released here
```



- Multiple readers OR single writer
- Good for read-heavy workloads
- Higher overhead than Mutex

```
use std::sync::RwLock;

let data = RwLock::new(vec![1, 2, 3]);

// Multiple readers
let r1 = data.read().unwrap();
let r2 = data.read().unwrap();

// Only one writer
let mut w = data.write().unwrap();
w.push(4);
```



- Lock ordering matters
- Try to acquire locks in same order
- Use scope-based RAII locks, shortest possible lifetime
- Nested lock problems (use ReentrantMutex from parking_lot)

```
let mutex1 = Mutex::new(0);
```

```
let mutex2 = Mutex::new(0);
```

```
// Potential deadlock!
```

```
let _guard1 = mutex1.lock().unwrap();
```

```
let _guard2 = mutex2.lock().unwrap();
```

```
// Better: always lock in same order
```

```
// across all code paths
```



- Wait for condition to become true
- Avoids busy waiting
- Always paired with mutex
- Allows thread park/wake



```
use std::sync::{Arc, Mutex, Condvar};
use std::thread;

let pair = Arc::new((Mutex::new(false), Condvar::new()));
let pair2 = Arc::clone(&pair);

thread::spawn(move || {
    let (lock, cvar) = &*pair2;
    let mut started = lock.lock().unwrap();
    *started = true;
    cvar.notify_one();
});
let (lock, cvar) = &*pair;
let mut started = lock.lock().unwrap();
while !*started {
    started = cvar.wait(started).unwrap();
}
```

Memory Models and Atomics



- Rust uses C++20 memory model
- Pragmatic choice for tooling
- Based on happens-before relationships
- Bridges hardware and software needs



- Compilers optimize aggressively
- May change execution order
- Example transformation:

```
// Original code
```

```
x = 1;
```

```
y = 3;
```

```
x = 2;
```

```
// What compiler might do
```

```
x = 2;
```

```
y = 3;
```




- x86/64: Strongly ordered
 - Most operations implicitly synchronized
 - Cheaper to provide strong guarantees
 - Might hide incorrect sync code
- ARM/RISC: Weakly ordered
 - Operations freely reordered
 - Explicit barriers needed
 - Better tests concurrent code



- Data accesses
 - Unsynchronized
 - Can be reordered
 - Cause data races
 - Cannot be used for synchronization
- Atomic accesses
 - Thread-aware
 - Ordering guarantees
 - Safe for synchronization
 - Various strength levels



- SeqCst (Strongest)
 - Global operation order
 - All threads agree
 - Most expensive
 - Safe default choice
- Release-Acquire
 - Paired operations
 - Good for locks
 - Prior writes visible
 - Cheaper than SeqCst
- Relaxed (Weakest)
 - Only atomicity
 - No synchronization
 - Good for counters
 - Best performance



```
use std::sync::atomic::{AtomicBool, Ordering};

static FLAG: AtomicBool = AtomicBool::new(false);

// All threads will agree on operation order
FLAG.store(true, Ordering::SeqCst);
let x = FLAG.load(Ordering::SeqCst);
```



```
use std::sync::atomic::{AtomicBool, Ordering};

static LOCK: AtomicBool = AtomicBool::new(false);

// Thread 1: Release ensures all writes visible
LOCK.store(false, Ordering::Release);

// Thread 2: Acquire sees released writes
while LOCK.compare_exchange(
    false, true,
    Ordering::Acquire,
    Ordering::Relaxed
).is_err() {}
```



```
use std::sync::atomic::{AtomicUsize, Ordering};

static COUNTER: AtomicUsize = AtomicUsize::new(0);

// Simple counter, no synchronization needed
thread::spawn(move || {
    COUNTER.fetch_add(1, Ordering::Relaxed);
});
```



- CPU instruction enforcing memory order
- Prevents reordering across fence
- Types:
 - LoadLoad: Orders loads before/after
 - StoreStore: Orders stores before/after
 - LoadStore: Orders loads before stores
 - StoreLoad: Strongest, orders all access

```
use std::sync::atomic::fence;
```

```
use std::sync::atomic::Ordering;
```

```
// Ensure all previous stores visible  
fence(Ordering::Release);
```

```
// Strongest fence, ensures all ordering  
fence(Ordering::SeqCst);
```

Message Passing



- MPSC = Multiple Producer, Single Consumer
- Thread-safe message queue
- Send values between threads
- Built into standard library

```
use std::sync::mpsc::channel;
```

```
let (tx, rx) = channel();
```

```
thread::spawn(move || {  
    tx.send(42).unwrap();  
});
```

```
println!("Got: {}", rx.recv().unwrap());
```



- `channel()` - Unbounded sync
- `sync_channel(n)` - Bounded sync
- Bounded channels block when full
- Choose based on backpressure needs
- Many types of channels (spsc, mpsc, etc.)

```
use std::sync::mpsc::{sync_channel, channel};
```

```
// Unbounded
```

```
let (tx1, rx1) = channel();
```

```
// Bounded to 5 messages
```

```
let (tx2, rx2) = sync_channel(5);
```

```
// Will block if buffer full
```

```
tx2.send(42).unwrap();
```



Flume

- Drop-in replacement for `std::sync::mpsc`
- Better performance
- Bounded/unbounded variants
- Select operation support

```
use flume::unbounded;
```

```
let (tx, rx) = unbounded();
```

```
thread::spawn(move || {  
    tx.send("Hello from flume!").unwrap();  
});
```

```
assert_eq!(rx.recv().unwrap(), "Hello from flume!");
```

Real World Patterns



- Common concurrent pattern
- Multiple producers feed queue
- Single consumer processes items

```
let (tx, rx) = flume::unbounded();  
// Spawn multiple producers  
for i in 0..3 {  
    let tx = tx.clone();  
    thread::spawn(move || {  
        tx.send(format!("msg from {}", i)).unwrap();  
    });  
}  
// Single consumer  
for msg in rx.iter() {  
    println!("Got: {}", msg);  
}
```



- Dynamic task distribution
- Threads steal work when idle
- Good for variable workloads

```
use crossbeam_deque::{Worker, Stealer};
let w = Worker::new_fifo();
let s = w.stealer();
// Main thread produces work
w.push(42);
// Other thread steals work
thread::spawn(move || {
    while let Some(work) = s.steal() {
        println!("Stolen: {}", work);
    }
});
```



- Reuse threads instead of creating new
- Better performance
- Control system resource usage
- Common in real applications

```
use rayon::prelude::*;
```

```
let data: Vec<_> = (0..100).collect();
```

```
// Parallel iterator using thread pool
```

```
let sum: i32 = data.par_iter()  
    .map(|x| x * x)  
    .sum();
```

Advanced Patterns



- Atomic operations have ordering guarantees
- Different levels of synchronization
- Performance vs. strictness tradeoff

```
use std::sync::atomic::{AtomicBool, Ordering};

static STOP: AtomicBool = AtomicBool::new(false);

// Relaxed - fastest, weakest guarantees
STOP.store(true, Ordering::Relaxed);

// Release-Acquire - synchronizes data
STOP.store(true, Ordering::Release);
let must_stop = STOP.load(Ordering::Acquire);

// SeqCst - strongest, slowest
STOP.store(true, Ordering::SeqCst);
```



- Coordinate multiple threads
- Wait for all threads to reach point
- Useful for phased computations

```
use std::sync::{Arc, Barrier};
```

```
let barrier = Arc::new(Barrier::new(3));
```

```
for _ in 0..3 {  
    let b = Arc::clone(&barrier);  
    thread::spawn(move || {  
        println!("before barrier");  
        b.wait(); // Wait for others  
        println!("after barrier");  
    });  
}
```



- No mutex required
- Uses atomic operations
- Higher performance potential
- More complex to implement

```
use crossbeam_queue::ArrayQueue;
```

```
let queue = ArrayQueue::new(100);
```

```
// Producer
```

```
queue.push(42).unwrap();
```

```
// Consumer
```

```
while let Some(item) = queue.pop() {  
    println!("Got: {}", item);  
}
```



- Implement Send and Sync traits
- Make your types thread-safe
- Careful with unsafe code

```
struct MyThreadSafeType {  
    data: Arc<Mutex<Vec<i32>>>,  
}  
  
// Automatically Send + Sync because  
// all fields are Send + Sync  
impl MyThreadSafeType {  
    fn push(&self, value: i32) {  
        self.data.lock().unwrap().push(value);  
    }  
}
```

Best Practices



- Channels for message passing
- Mutex for shared state
- Atomic for simple values
- Thread pools for CPU work

```
// Bad: Mutex for single counter
```

```
let counter = Arc::new(Mutex::new(0));
```

```
// Better: Atomic for counter
```

```
let counter = Arc::new(AtomicUsize::new(0));
```



- Measure before optimizing
- Consider contention
- Right granularity of locks
- Cache coherency effects

```
// Bad: Too fine-grained
let numbers: Vec<_> = (0..1000)
    .map(|i| Arc::new(Mutex::new(i)))
    .collect();
```

```
// Better: Coarser granularity
let numbers = Arc::new(Mutex::new(
    (0..1000).collect::<Vec<_>>()
));
```

Questions?
